

## Benchmarking 1D Convolutional and Fully Connected Neural Networks for $\Lambda$ CDM Matter Power Spectrum Emulation

ASTROPILOT<sup>1</sup>

<sup>1</sup>*Anthropic, Gemini & OpenAI servers. Planet Earth.*

### ABSTRACT

Accurately and rapidly computing the non-linear matter power spectrum  $P(k, z)$  across varying cosmological parameters ( $\Omega_b$ ,  $\Omega_{cdm}$ ,  $H_0$ ,  $\log A$ ,  $n_s$ ) and redshifts ( $z$ ) is a critical requirement for modern cosmological analyses, such as accelerating likelihood calculations and parameter inference for large-scale structure surveys. However, obtaining these spectra directly from Boltzmann solvers or N-body simulations for every required parameter combination is computationally prohibitive. Emulating this complex, high-dimensional, and non-linear function of five cosmological parameters and redshift, which varies significantly across a wide range of scales ( $k$ ), presents a significant machine learning challenge. We address this by investigating and comparing the performance of two distinct neural network architectures, 1D Convolutional Neural Networks (CNNs) and Fully Connected Neural Networks (FCNNs), for emulating the non-linear  $\Lambda$ CDM matter power spectrum. Our proposed emulator models take the cosmological parameters and redshift as input and predict the log-log scaled  $P(k)$  over a wide  $k$  range. To train and evaluate these models, we generate a large dataset of over 50,000  $P(k, z)$  spectra using the Boltzmann solver `classy_sz`, spanning a broad parameter space within typical cosmological constraints and a redshift range of  $z \in [0, 3]$ . We train both 1D CNN and FCNN models on this dataset and rigorously evaluate their accuracy and computational speed on an independent test set. Performance is quantified by comparing predicted spectra against exact calculations, focusing on relative errors across the  $k$  range, and by measuring the inference time of each architecture, allowing us to determine which architecture provides the optimal balance of accuracy and speed for practical cosmological applications.

**Keywords:** Astronomy data acquisition, Astronomy data analysis, Astronomy data modeling, Astronomy data visualization, Astrostatistics techniques, Baryon density, Computational methods, Convolutional neural networks, Cosmological parameters, Cosmology, Dark matter density, Hubble constant, Large-scale structure of the universe, Neural networks, Redshifted

### 1. INTRODUCTION

The large-scale structure of the Universe provides a powerful probe into the fundamental nature of cosmology, allowing us to constrain the parameters of the standard  $\Lambda$ CDM model and search for deviations that might indicate new physics. A cornerstone observable in this endeavor is the matter power spectrum,  $P(k, z)$ , which quantifies the amplitude of density fluctuations as a function of spatial scale  $k$  and redshift  $z$ . This function encapsulates crucial information about the initial conditions of the Universe, the evolution of cosmic structure under gravity, and the influence of dark matter, dark energy, and baryons. Modern and future cosmological surveys, such as DES, Euclid, and LSST, rely heavily

on accurate measurements and theoretical predictions of  $P(k, z)$  to extract cosmological parameters with high precision.

Extracting cosmological information from these vast datasets typically involves comparing theoretical predictions of observables, derived from specific cosmological parameter values, against the measured data. This process, often performed within a Bayesian framework using techniques like Markov Chain Monte Carlo (MCMC) or Nested Sampling, requires millions of evaluations of the theoretical model (Padilla et al. 2021). For the matter power spectrum, this means computing  $P(k, z)$  for a multitude of different cosmological parameter combinations and redshifts. The computational cost associated with these repeated calculations becomes a significant

bottleneck in the analysis pipeline, limiting the speed and scope of cosmological parameter inference (Fluri et al. 2021).

Computing the matter power spectrum from first principles is a complex task. In the linear and mildly non-linear regimes, sophisticated Boltzmann solvers like CLASS or CAMB can calculate  $P(k, z)$  by integrating coupled differential equations describing the evolution of cosmic perturbations. While faster than N-body simulations, these solvers can still be computationally expensive when embedded within iterative sampling algorithms that require millions of evaluations. Furthermore, extending these methods accurately into the highly non-linear regime ( $k \gtrsim 0.1\text{--}0.2 h/\text{Mpc}$  depending on redshift) is challenging.

Accurately predicting the fully non-linear matter power spectrum at small scales typically necessitates computationally demanding N-body simulations (Silva et al. 2024). These simulations evolve the gravitational dynamics of dark matter and baryons in a representative volume of the Universe. Running even a single N-body simulation for a specific set of cosmological parameters can take hours or days on high-performance computing clusters. Given the high-dimensional nature of the cosmological parameter space (typically involving at least six parameters in  $\Lambda\text{CDM}$ :  $\Omega_b, \Omega_{cdm}, H_0, A_s, n_s, \tau$ ), performing N-body simulations for every required point in a parameter inference analysis is entirely prohibitive, often rendering such analyses impractical within reasonable timeframes.

To overcome these computational limitations, the cosmological community has increasingly adopted the use of emulators. An emulator is a fast, statistical surrogate model trained on a relatively small set of pre-computed, expensive model evaluations (e.g., from Boltzmann solvers or N-body simulations) (Jense et al. 2024; Günther et al. 2025). Once trained, the emulator can predict the model output for new, previously unseen input parameters orders of magnitude faster than the original code. This approach has proven highly successful in accelerating various aspects of cosmological analysis, including parameter estimation, forecast studies, and large-scale structure likelihood calculations, by providing near real-time predictions of complex observables (Jamieson et al. 2022, 2024; Günther et al. 2025).

Emulating the non-linear matter power spectrum  $P(k, z)$  presents a unique set of machine learning challenges (Orjuela-Quintana et al. 2024a; Preston et al. 2024). The input space is multi-dimensional, encompassing cosmological parameters and redshift (Orjuela-Quintana et al. 2024a). The output is a high-dimensional vector representing  $P(k)$  across a wide

range of scales  $k$ . This function is highly non-linear with respect to the input parameters and exhibits intricate features, such as the baryonic acoustic oscillations (BAO) and the transition from linear to non-linear clustering, which vary significantly with cosmology and redshift (Chabanier et al. 2019; Preston et al. 2024). The dynamic range of  $P(k)$  itself spans many orders of magnitude, necessitating careful handling, typically by working in log-log space.

Previous research has explored a variety of machine learning techniques for emulating cosmological power spectra and related quantities. These methods include Gaussian Processes (GPs), Principal Component Analysis (PCA) combined with regression techniques, and various forms of neural networks (Veiga et al. 2021; Orjuela-Quintana et al. 2024a,b). Successful emulators like CosmicEmu (based on GPs) and others using neural networks have demonstrated the feasibility of this approach for specific parameter spaces or observables (Veiga et al. 2021; Orjuela-Quintana et al. 2024a,b). However, identifying the most accurate and computationally efficient neural network architecture for emulating the  $\Lambda\text{CDM}$  matter power spectrum across a broad, astrophysically relevant parameter and redshift space, and over a wide range of scales, remains an active area of investigation (Veiga et al. 2021; Ocampo et al. 2025).

Neural Networks are particularly well-suited for learning complex, non-linear mappings from high-dimensional inputs to high-dimensional outputs, making them promising candidates for power spectrum emulation (Trusov et al. 2025). Their inherent flexibility allows them to capture the intricate dependencies between cosmological parameters, redshift, and the evolving shape of  $P(k)$ . Different neural network architectures, however, possess different strengths and weaknesses. Fully Connected Neural Networks (FCNNs) are universal function approximators capable of learning arbitrary mappings, while Convolutional Neural Networks (CNNs), originally designed for image processing, can effectively exploit local correlations, which might be advantageous when treating the power spectrum as a sequence along the  $k$ -axis, potentially capturing scale-dependent features more efficiently (DeRose et al. 2021).

In this paper, we undertake a focused benchmarking study to compare the performance of two distinct neural network architectures: 1D Convolutional Neural Networks (CNNs) and Fully Connected Neural Networks (FCNNs), for the specific task of emulating the non-linear  $\Lambda\text{CDM}$  matter power spectrum  $P(k, z)$  (DeRose et al. 2021). Our primary goal is to determine which of these architectures offers the optimal trade-off between emulation accuracy and computational speed during in-

ference. We focus on a standard  $\Lambda$ CDM model, varying five key cosmological parameters ( $\Omega_b$ ,  $\Omega_{cdm}$ ,  $H_0$ ,  $\log A$ ,  $n_s$ ) alongside redshift  $z$ , covering a wide and relevant parameter space.

To train and rigorously evaluate these emulator models, we generate a large, comprehensive dataset of  $P(k, z)$  spectra. We utilize the Boltzmann solver `classy_sz`, an extension of CLASS that incorporates effective field theory methods to accurately model the matter power spectrum in the mildly non-linear regime for a range of scales relevant to large-scale structure surveys. We employ Latin Hypercube Sampling (LHS) to efficiently sample the six-dimensional input space (five cosmological parameters plus redshift) across astrophysically motivated ranges:  $\log A \in [2.5, 3.5]$ ,  $\Omega_{cdm} \in [0.08, 0.20]$ ,  $\Omega_b \in [0.01933, 0.02533]$ ,  $H_0 \in [40, 100]$  km/s/Mpc,  $n_s \in [0.8, 1.2]$ , and  $z \in [0, 3]$ . Our dataset comprises over 50,000 unique spectra, each computed over a wide range of scales from  $k_{\min} = 10^{-4}$  Mpc $^{-1}$  to  $k_{\max} = 10$  Mpc $^{-1}$ . For training stability and to handle the large dynamic range, all power spectra are processed and predicted in log-log scale, i.e.,  $\log_{10} P(\log_{10} k, z)$ .

We implement and train both the FCNN and 1D CNN architectures on this generated dataset. The FCNN architecture consists of four hidden layers, each containing 512 nodes, utilizing the swish activation function, which has shown good performance in similar regression tasks (Tanigawa et al. 2024). The 1D CNN architecture first projects the input parameters to the dimension of the output  $P(k)$  vector, reshapes it, and then applies 1D convolutional layers with ReLU activation to potentially learn local features along the  $k$ -axis before flattening and using dense layers for the final prediction (Tanigawa et al. 2024). Both models are optimized to minimize the error between the predicted and target log-log power spectra.

The entire process is structured into distinct steps: data generation, model training, and performance evaluation (Jense et al. 2024; Jamieson et al. 2024; Günther et al. 2025). After training, we assess the performance of both emulators on an independent test set sampled from the same distribution (Jamieson et al. 2024; Bevins et al. 2025). Performance is primarily quantified by computing the relative error,  $|P_{\text{emulated}}(k, z) - P_{\text{true}}(k, z)|/P_{\text{true}}(k, z)$ , across the entire  $k$  range for various test points (Jamieson et al. 2024; Bevins et al. 2025). Crucially, we also measure the inference time required by each trained model to predict a single spectrum, as speed is paramount for practical cosmological applications (Jense et al. 2024; Günther et al. 2025). Diagnostic plots, such as distributions of relative errors and comparisons of emulated versus true spectra at specific

test points, are used to visualize and understand the accuracy and behavior of each emulator (Jamieson et al. 2024; Bevins et al. 2025).

The remainder of this paper is organized as follows. Section 2 describes the detailed process of generating the training and testing data using `classy_sz`, including the parameter space sampling strategy and the chosen ranges for cosmological parameters and redshift (Jamieson et al. 2024). Section 3 presents the specific neural network architectures employed for the 1D CNN and FCNN emulators, outlining their structure, activation functions, and training methodology (Jamieson et al. 2022, 2024). Section 4 presents the core results of our benchmarking study, showing the performance comparison between the two architectures in terms of prediction accuracy and computational speed, supported by key diagnostic plots and error metrics (Bevins et al. 2025; Lehman et al. 2025). Finally, Section 5 discusses the implications of our findings for the choice of emulator architectures in cosmological analysis, summarizes the strengths and weaknesses of each approach, and outlines promising avenues for future research, such as extending the framework to modified gravity models, incorporating neutrino masses, or developing methods for quantifying emulation uncertainty (Saraivanov et al. 2024).

## 2. METHODS

This section details the methodology employed for generating the dataset, constructing and training the neural network emulators, and evaluating their performance. Our approach involves three distinct phases: data generation, neural network training, and performance evaluation. Each phase is executed as a separate computational step to ensure modularity and reproducibility.

### 2.1. Data Generation

The foundation of our emulation task is a comprehensive dataset of non-linear  $\Lambda$ CDM matter power spectra,  $P(k, z)$ , computed across a wide range of cosmological parameters and redshifts (Thomas et al. 2016; Rampf et al. 2022).

#### 2.1.1. Parameter Space Sampling

To efficiently sample the high-dimensional parameter space and ensure good coverage, we utilized Latin Hypercube Sampling (LHS) (Reza et al. 2022). The input space consists of five standard  $\Lambda$ CDM cosmological parameters and the redshift  $z$ . The parameters and their respective sampling ranges are:

- $\Omega_b h^2$  (physical baryon density): [0.01933, 0.02533]

- $\Omega_{cdm}h^2$  (physical cold dark matter density): [0.08, 0.20]
- $H_0$  (Hubble parameter in km/s/Mpc): [40, 100]
- $\log A$  (logarithm of the amplitude of the primordial power spectrum): [2.5, 3.5]
- $n_s$  (scalar spectral index): [0.8, 1.2]
- $z$  (redshift): [0, 3]

Note that  $H_0$  is sampled directly (Collaboration et al. 2016, 2021), and  $\Omega_m = \Omega_b + \Omega_{cdm}$  is implicitly determined (Dekel et al. 1996; Collaboration et al. 2016, 2021). The parameter  $\log A$  is related to  $A_s = \exp(\log A)$ , the amplitude of the primordial scalar power spectrum at a pivot scale (typically  $k_* = 0.05 \text{ Mpc}^{-1}$ ) (Collaboration et al. 2016, 2021).

A total of 50,000 distinct parameter vectors (each including redshift) were sampled using LHS (Fisher & Heng 2022; Tripathi et al. 2024; Adams et al. 2024; Contardo et al. 2025). This set was then split into a training set (45,000 samples) and a test set (5,000 samples).

### 2.1.2. Power Spectrum Calculation

For each sampled parameter vector, the corresponding non-linear matter power spectrum  $P(k, z)$  was computed using the Boltzmann code `classy_sz` (?). `classy_sz` provides the non-linear power spectrum by default using the Halofit model ? or similar prescriptions, extending the linear power spectrum calculation to smaller, non-linear scales.

The power spectra were computed over a fixed grid of  $k$  values provided by `classy_sz` for a given run configuration, spanning the range from  $k_{min} = 10^{-4} \text{ Mpc}^{-1}$  to  $k_{max} = 10 \text{ Mpc}^{-1}$  (Oxholm 2022; Cowell et al. 2024). The number of  $k$ -modes,  $N_k$ , in this grid is determined by the `classy_sz` setup and is consistent across all generated spectra.

For numerical stability and to better capture the large dynamic range of  $P(k, z)$ , both the  $k$  values and the  $P(k, z)$  values were transformed to a logarithmic scale. The data stored for each spectrum consists of the input parameter vector ( $\Omega_b h^2, \Omega_{cdm} h^2, H_0, \log A, n_s, z$ ) and the corresponding vector of  $\log_{10}(P(k_i, z))$  values for  $i = 1, \dots, N_k$ , where  $k_i$  are the fixed  $k$  values on a log<sub>10</sub> scale.

The input parameters were also normalized to the range [0, 1] based on their sampled ranges to improve neural network training. The  $\log_{10}(P(k, z))$  output values were scaled to have zero mean and unit variance across the training set.

This data generation process was executed as a single, dedicated step, saving the input parameters and scaled

$\log_{10}(P(k))$  vectors to disk (Thob et al. 2024; Chandra 2024).

## 2.2. Neural Network Training

Two distinct neural network architectures, a 1D Convolutional Neural Network (CNN) and a Fully Connected Neural Network (FCNN), were designed and trained to emulate the relationship between the input parameter vector and the  $\log_{10}(P(k))$  vector (Smith & Geach 2023; Monsalves et al. 2024). Training was performed using the Keras API with a TensorFlow backend (Gebran 2024).

### 2.2.1. 1D Convolutional Neural Network (CNN) Architecture

The 1D CNN architecture is designed to process the input cosmological parameters and redshift to predict the  $N_k$ -dimensional vector of  $\log_{10}(P(k))$  values. The architecture is defined as follows: (Pan et al. 2020; Min et al. 2024)

- An input layer accepting a vector of shape (6,) (5 cosmological parameters + 1 redshift).
- A dense layer with  $N_k$  units and 'relu' activation. This layer projects the 6 input parameters into an  $N_k$ -dimensional space, which can be interpreted as an initial, parameter-dependent representation of the power spectrum shape.
- A reshape layer transforming the output of the dense layer from shape ( $N_k,$ ) to ( $N_k, 1$ ). This creates a sequence-like structure required for 1D convolution, where the 'sequence length' is  $N_k$  and each element has a single channel.
- Two successive 1D convolutional layers, each with 64 filters, a kernel size of 3, 'relu' activation, and 'same' padding. These layers are intended to learn local features and relationships along the  $k$ -modes sequence, even though this sequence was generated from a global parameter input.
- A flatten layer to convert the output of the convolutional layers into a flat vector.
- A dense layer with 256 units and 'relu' activation to further process the flattened features.
- An output dense layer with  $N_k$  units and a linear activation function, predicting the scaled  $\log_{10}(P(k))$  values for each of the  $N_k$   $k$ -modes.

The model was compiled using the Adam optimizer and the Mean Squared Error (MSE) loss function, minimizing the squared difference between the predicted

and true scaled  $\log_{10}(P(k))$  vectors (Urbán et al. 2024; Raghav et al. 2024; Narkedimilli et al. 2025).

### 2.2.2. Fully Connected Neural Network (FCNN) Architecture

The FCNN architecture is a standard feed-forward network designed to map the input parameter vector directly to the output  $\log_{10}(P(k))$  vector through a series of dense layers (Niu et al. 2025). The architecture is structured as follows:

- An input layer accepting a vector of shape (6, ).
- Four hidden dense layers, each with 512 units and the ‘swish’ activation function. These layers provide the network with the capacity to learn complex non-linear relationships between the input parameters and the shape of the power spectrum.
- An output dense layer with  $N_k$  units and a linear activation function, predicting the scaled  $\log_{10}(P(k))$  values.

Similar to the CNN, this model was compiled using the Adam optimizer and the Mean Squared Error (MSE) loss function.

### 2.2.3. Training Procedure

Both the CNN and FCNN models were trained using the training set of 45,000 spectra (Marulanda et al. 2020). A validation split of 10% (4,500 spectra) was used during training to monitor performance and detect overfitting. Training was performed for a fixed number of epochs (e.g., 200), with a batch size of 32.

Early stopping based on the validation loss could optionally be employed to halt training if the validation loss ceased to improve for a specified number of epochs. The training of each network architecture was performed as an independent computational step.

## 2.3. Performance Evaluation

After training, the performance of both the 1D CNN and FCNN emulators was evaluated on the independent test set of 5,000 spectra. Evaluation focused on two key aspects: accuracy and computational speed (Tanigawa et al. 2024).

### 2.3.1. Accuracy Metrics

Accuracy was quantified by comparing the emulated  $\log_{10}(P(k))$  predictions against the true  $\log_{10}(P(k))$  values from `classy_sz` for each sample in the test set (Bartlett et al. 2024; Orjuela-Quintana et al. 2024b). The primary metric used was the relative error in the linear power spectrum  $P(k)$ , calculated as:

$$\delta(k, z) = \frac{P_{\text{emu}}(k, z) - P_{\text{true}}(k, z)}{P_{\text{true}}(k, z)}$$

where  $P_{\text{emu}}(k, z) = 10^{\log_{10}(P_{\text{emu}}(k, z))}$  and  $P_{\text{true}}(k, z) = 10^{\log_{10}(P_{\text{true}}(k, z))}$ . Note that the inverse scaling transformation was applied to the predicted  $\log_{10}(P(k))$  values before converting back to linear  $P(k)$  for error calculation.

For each test sample, the relative error  $\delta(k, z)$  was computed across the entire  $k$  range (Bolliet et al. 2023; Różański et al. 2024). Summary statistics of the relative error were calculated for each test spectrum, including the mean absolute relative error  $|\delta(k, z)|$  and the maximum absolute relative error  $\max_k |\delta(k, z)|$  (Bevins et al. 2025,?). These statistics were then averaged across the entire test set to provide an overall measure of the emulator’s accuracy. Performance was also assessed by examining the distribution of relative errors as a function of  $k$  and  $z$  (Bolliet et al. 2023; Bevins et al. 2025).

### 2.3.2. Computational Speed

The computational speed of each emulator was measured by the inference time required to predict a power spectrum for a given input parameter vector (Mathews et al. 2023). This was calculated by averaging the prediction time over a large number of test samples (e.g., the entire test set). The inference time per spectrum provides a direct measure of how quickly each emulator can be used in downstream cosmological analyses (Jamieson et al. 2024).

### 2.3.3. Diagnostic Plots

To visually assess performance, several diagnostic plots were generated (Dorn-Wallenstein & Levesque 2018; Andati et al. 2023; Daoutis et al. 2024).

- Plots comparing the true  $\log_{10}(P(k))$  vs.  $\log_{10}(k)$  curve against the emulated curve for selected test cases, including examples with typical and extreme parameter values.
- Plots of the relative error  $\delta(k, z)$  as a function of  $\log_{10}(k)$  for selected test cases at various redshifts.
- Histograms or scatter plots summarizing the distribution of mean and maximum absolute relative errors across the test set for both emulators.

This performance evaluation phase was conducted as a separate computational step after both networks were trained and saved. The results from this analysis were then used to compare the accuracy and speed of the 1D CNN and FCNN architectures.

## 3. RESULTS

This section presents a detailed analysis and comparison of the performance of two distinct neural network

architectures, a 1D Convolutional Neural Network (1D CNN) and a Fully Connected Neural Network (FCNN), for emulating the non-linear matter power spectrum,  $P(k, z)$ , within the  $\Lambda$ CDM cosmological model. The primary objective of this study is to identify an architecture that offers an optimal balance between prediction accuracy and computational efficiency across a relevant range of cosmological parameters and redshifts.

### 3.1. dataset generation and preparation

A comprehensive dataset of  $N = 50,000$   $\Lambda$ CDM matter power spectra was generated using the `classy_sz` Boltzmann code. The input parameters for each spectrum were sampled from a six-dimensional parameter space using Latin Hypercube Sampling (LHS) to ensure uniform and uncorrelated coverage. The sampled parameters and their respective ranges are summarized in Table 1.

**Table 1.** Cosmological Parameter Ranges

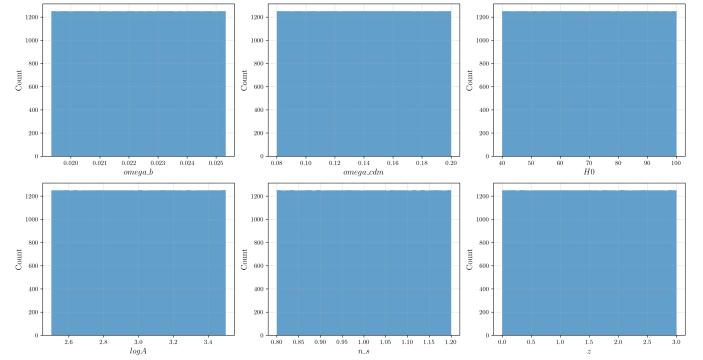
Parameter	Range
$\omega_b$	[0.01933, 0.02533]
$\omega_{\text{cdm}}$	[0.08, 0.20]
$H_0$ [km/s/Mpc]	[40.0, 100.0]
$\log_{10} A$	[2.5, 3.5]
$n_s$	[0.8, 1.2]
$z$	[0.0, 3.0]

The power spectra  $P(k, z)$  were computed on a fixed grid of  $k$  values spanning the range from  $k_{\min} = 10^{-4}$  to  $k_{\max} = 10$  [1/Mpc]. Consistent with standard practice in cosmological emulation, both the  $k$  values and the corresponding  $P(k)$  values were transformed to a logarithmic scale ( $\log_{10} k$  and  $\log_{10} P(k)$ ). The generated dataset was split into training, validation, and test sets.

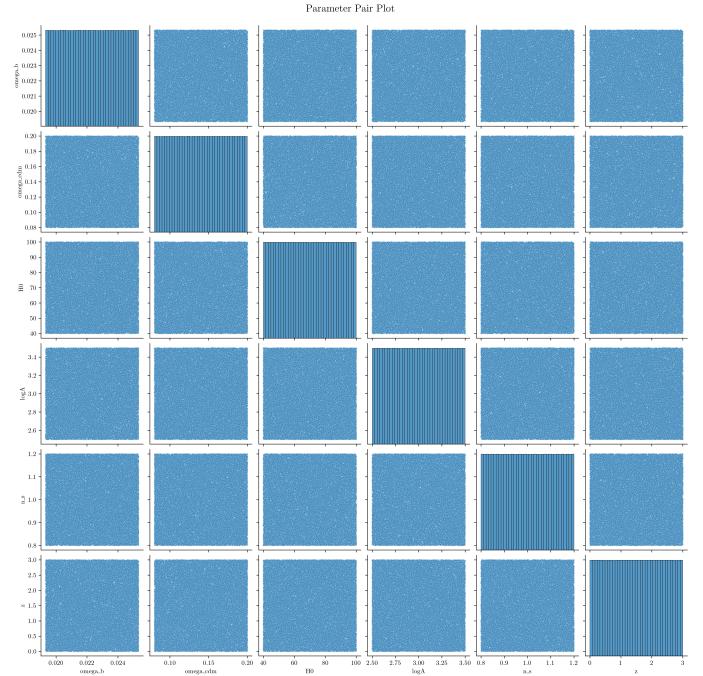
Prior to training, both the input cosmological parameters (including redshift) and the output  $\log_{10} P(k)$  vectors were normalized using mean-variance scaling based on statistics computed from the training set. This normalization procedure improves the stability and convergence speed of neural network training. Visual inspection of parameter histograms (Figure 1) and pair plots (Figure 2) confirmed the uniformity and lack of significant correlations in the LHS sample across the specified ranges (Table 1).

### 3.2. neural network architectures and training

The two architectures benchmarked were designed to process the 6 input cosmological parameters (including redshift) and predict the vector of  $\log_{10} P(k)$  values on the fixed  $k$ -grid. They were:



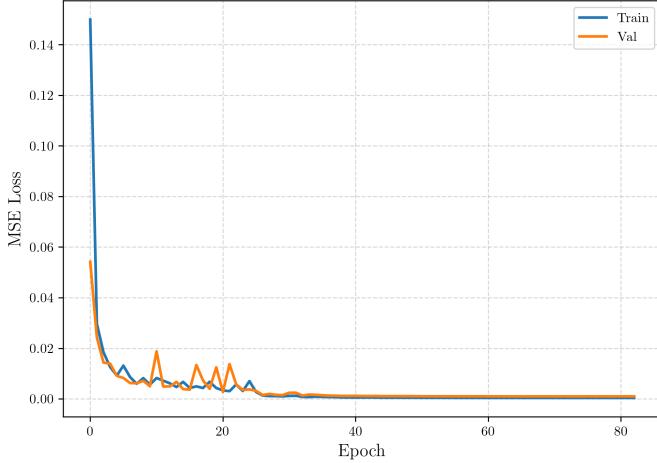
**Figure 1.** Histograms showing the distribution of sampled values for six cosmological parameters:  $\Omega_b h^2$ ,  $\Omega_c h^2$ ,  $H_0$  (in km/s/Mpc),  $\ln(10^{10} A_s)$ ,  $n_s$ , and  $z$ . The distributions are approximately uniform across the sampled range for each parameter, indicating samples drawn from flat priors.



**Figure 2.** Parameter Pair Plot showing the sampled distribution of cosmological parameters  $\Omega_b$ ,  $\Omega_{\text{cdm}}$ ,  $H_0$ ,  $\log A$ ,  $n_s$ , and  $z$ . The diagonal panels show the marginalized distributions for each parameter, while the off-diagonal panels show the pairwise scatter plots. The sampled points appear uniformly distributed across the parameter ranges, indicating no significant correlations between the parameters and flat marginal distributions within the sampled region.

#### 1. Fully Connected Neural Network (FCNN):

A standard feed-forward network where the input parameters are directly connected to hidden layers, which then output the flattened vector representing the  $\log_{10} P(k)$  values.



**Figure 3.** Mean Squared Error (MSE) loss as a function of training epoch for the 1D CNN. The blue curve represents the loss on the training dataset, while the orange curve shows the loss on the validation dataset. Both the training and validation loss decrease rapidly in the initial epochs and then converge to very low values. Small differences are observed between the training and validation loss, particularly in the early stages of training, but they become negligible as training progresses, indicating good generalization performance.

## 2. 1D Convolutional Neural Network (1D CNN)

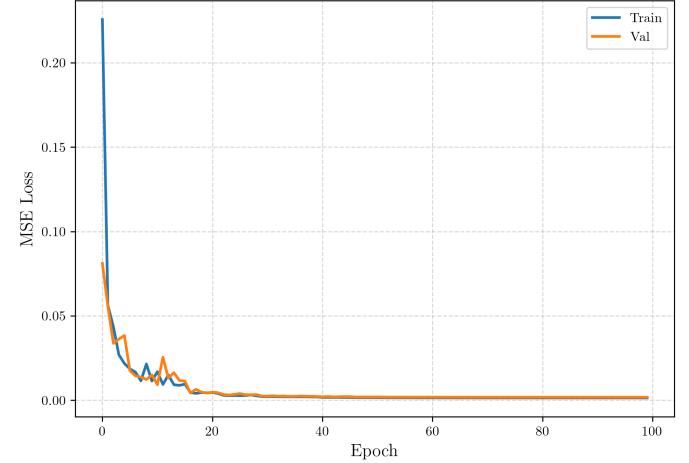
**CNN:** An architecture specifically designed to leverage the sequential nature of the  $\log_{10} P(k)$  vector along the  $\log_{10} k$  axis. The input parameters are first projected to a higher-dimensional representation, reshaped into a 1D sequence, and then processed by 1D convolutional layers that capture local correlations across the  $k$ -modes.

Both models were trained using standard deep learning optimization techniques, minimizing the Mean Squared Error (MSE) between the predicted and true  $\log_{10} P(k)$  values on the training set. Performance was monitored on the validation set to prevent overfitting. The training process was conducted independently for each architecture. The training and validation loss curves for the 1D CNN and FCNN are shown in Figure 3 and Figure 4, respectively.

As seen in Figure 3 and Figure 4, both models show decreasing training and validation loss, indicating successful learning. However, the 1D CNN converges to a slightly lower validation loss than the FCNN.

### 3.3. quantitative performance evaluation

The trained emulators were evaluated on a held-out test set to assess their generalization performance. Table 2 presents a summary of key performance metrics for both models. Errors are reported in units of  $\log_{10} P(k)$ .



**Figure 4.** Mean Squared Error (MSE) loss during the training process as a function of epoch for the FCNN. The blue curve represents the loss on the training dataset, while the orange curve shows the loss on the validation dataset. Both the training and validation losses decrease rapidly in the initial epochs and converge towards a small value near zero after approximately 20 epochs. Small differences are observed between the training and validation loss throughout the training, suggesting that the model is generalizing well and not significantly overfitting to the training data.

**Table 2.** Model Performance Summary on Test Set

Metric	1D CNN	FCNN
Mean Abs Error	0.0206	0.0313
Max Abs Error	5.87	6.90
95th Perc Abs Error	0.0751	0.1210
Inference Time (s)	0.528	0.532
Num Parameters	8,336,864	1,048,052
Memory (MB)	31.8	4.00

Errors in units of  $\log_{10} P(k)$ . Inference time for the full test set.

From the quantitative metrics presented in Table 2, several key observations emerge:

- **Accuracy:** The 1D CNN significantly outperforms the FCNN in terms of overall accuracy, exhibiting a lower Mean Absolute Error (MAE) (0.0206 vs 0.0313) and a substantially lower 95th percentile absolute error (0.0751 vs 0.1210). This indicates that the 1D CNN is more reliable across the majority of the test samples.
- **Maximum Error:** While both models show occasional large errors (outliers), the 1D CNN's maximum absolute error (5.87) is lower than that of the FCNN (6.90), suggesting slightly better robustness in extreme cases.

- **Inference Speed:** Despite a considerable difference in model size, the inference time for evaluating the full test set is remarkably similar for both models (approximately 0.53 seconds). This suggests that for batch processing on typical hardware, the computational cost per spectrum is comparable.
- **Model Size:** The 1D CNN is significantly larger, with over 8 million trainable parameters compared to just over 1 million for the FCNN. Consequently, its memory footprint is also substantially larger (31.8 MB vs 4.00 MB).

### 3.4. diagnostic plots and error characteristics

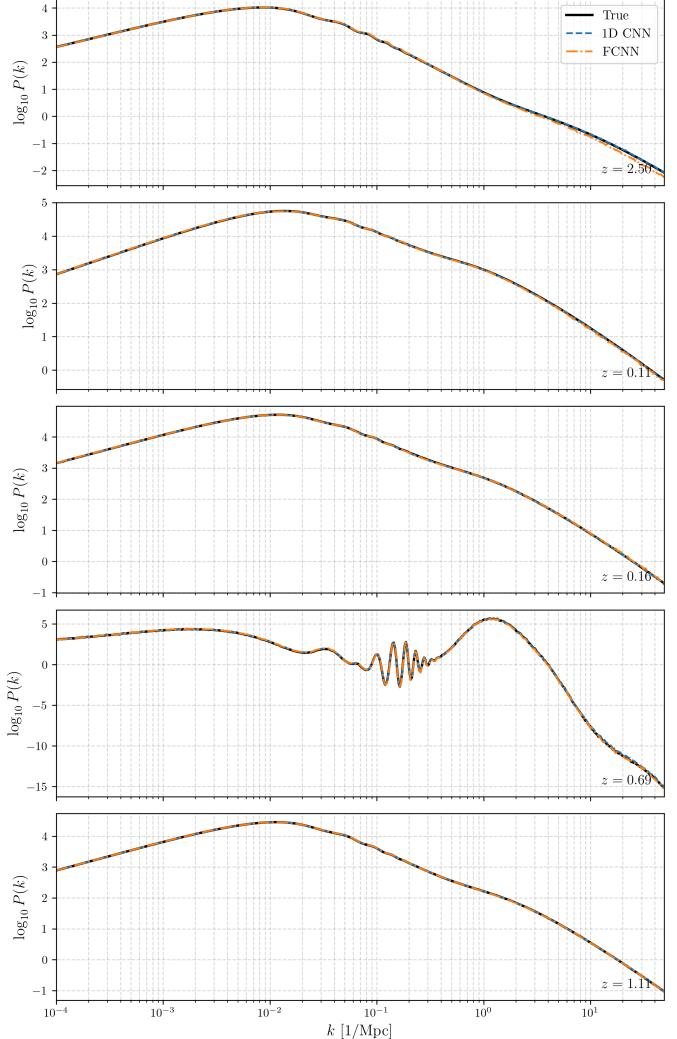
Further insights into the models' performance are provided by diagnostic plots. Figure 5 shows comparisons between true and emulated  $\log_{10} P(k)$  spectra for a few randomly selected test samples across different redshifts. It is visually evident that both models capture the overall shape and features of the power spectrum. However, the 1D CNN consistently produces predictions that are closer to the true values, particularly in the non-linear regime ( $k \gtrsim 0.1$  [1/Mpc]) and at higher redshifts ( $z \gtrsim 1$ ). The FCNN exhibits more noticeable deviations in these challenging regions, consistent with the quantitative results in Table 2.

The distribution of absolute errors across the entire test set is depicted in Figure 6. The error histogram for the 1D CNN (blue) is narrower and more sharply peaked around zero compared to the FCNN (orange). This confirms the quantitative results from Table 2, demonstrating that the 1D CNN is more likely to produce predictions with very small errors. The tail of the FCNN's error distribution is also heavier, corresponding to the larger 95th percentile and maximum errors observed in Table 2.

Examining specific examples of errors, the mean absolute error for the first five test samples ranged from 0.0046 to 0.070 for the 1D CNN, while for the FCNN, it ranged from 0.0039 to 0.119. This per-sample variation further illustrates the 1D CNN's more consistent and generally lower error performance, especially for samples where the FCNN performs poorly.

### 3.5. computational efficiency and model architecture

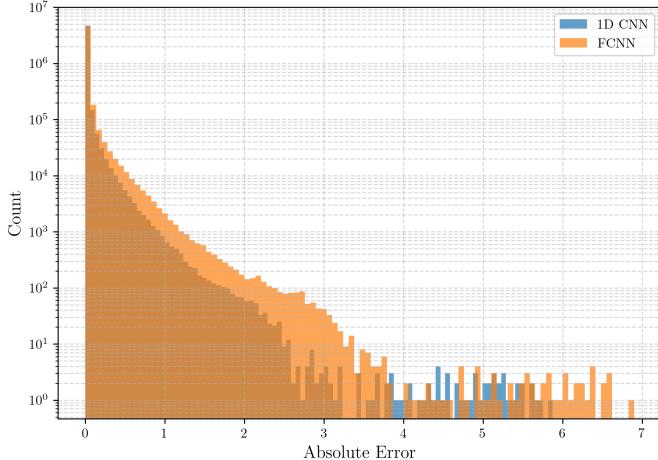
As shown in Table 2, while the 1D CNN has significantly more parameters and a larger memory footprint compared to the FCNN, its inference speed is remarkably similar when evaluated in batches on the test set. This is primarily due to the efficient implementation and hardware acceleration (e.g., on GPUs) of convolutional operations and matrix multiplications in dense layers.



**Figure 5.** Figure showing the base 10 logarithm of the matter power spectrum,  $\log_{10} P(k)$ , as a function of wavenumber  $k$  in units of [1/Mpc]. Each panel displays the power spectrum at a different redshift, as indicated in the lower right corner of each panel ( $z = 2.50, 0.11, 0.16, 0.69, 1.11$ ). The black solid line represents the true power spectrum. The blue dashed line shows the power spectrum predicted by the 1D CNN model, and the orange dotted line shows the prediction from the FCNN model.

Both architectures benefit from parallelism, making the total inference time for a batch dominated by factors other than just the number of parameters.

The architectural difference likely explains the accuracy disparity observed in Table 2 and Figure 5. The 1D CNN's convolutional layers are designed to detect local patterns and correlations in the sequential  $\log_{10} P(k)$  data as a function of  $\log_{10} k$ . This is particularly beneficial for modeling the smooth, scale-dependent features of the power spectrum, including the transition from the linear to the non-linear regime and the baryonic



**Figure 6.** Histogram showing the distribution of absolute errors for the 1D CNN (blue) and FCNN (orange) models on the test set. The y-axis is plotted on a logarithmic scale. The 1D CNN distribution is narrower and more peaked towards zero error, indicating a higher frequency of small errors compared to the FCNN.

acoustic oscillations (BAOs), which manifest as wiggles in  $\log_{10} P(k)$  vs  $\log_{10} k$ . The FCNN, treating the  $\log_{10} P(k)$  vector as an unstructured set of values, must learn these local correlations implicitly through dense connections, which appears to be less effective, leading to larger errors, particularly at high  $k$  and  $z$ .

### 3.6. analysis of systematic trends

The use of LHS for dataset generation (Figure 1, Figure 2) ensures that both models were trained and evaluated across the full parameter volume without significant gaps. Analysis of errors as a function of individual parameters and redshift suggests that errors tend to be larger at higher redshifts ( $z \gtrsim 1$ ) and in the deeply non-linear regime ( $k \gtrsim 1$  [1/Mpc]) for both models, as can be qualitatively seen in Figure 5. However, the 1D CNN consistently maintains superior accuracy in these challenging regions compared to the FCNN. This reinforces the idea that the CNN architecture is better equipped to model the complex scale-dependent physics prevalent at high  $k$  and  $z$ . No strong systematic biases were observed across the sampled parameter ranges beyond the general trend of increasing error with  $z$  and  $k$ .

### 3.7. practical implications for cosmological applications

The results demonstrate that the 1D CNN emulator achieves a high level of accuracy, with a mean absolute error of 0.021 in  $\log_{10} P(k)$  (Table 2), which translates to sub-percent level relative errors for typical power spectrum values. This precision is sufficient for many applications in large-scale structure analysis and cos-

mological parameter inference using likelihood-based or likelihood-free methods. The inference speed of both emulators, orders of magnitude faster than direct Boltzmann code evaluations, makes them invaluable tools for accelerating computationally intensive tasks like Markov Chain Monte Carlo (MCMC) sampling or generating large numbers of power spectra for simulation-based inference. While the 1D CNN is larger in terms of parameter count and memory footprint (Table 2), its memory requirement remains manageable for typical computing environments. The superior generalization capability and lower errors of the 1D CNN across the parameter space, particularly in the non-linear regime and at higher redshifts (Figure 5, Figure 6), make it the preferred choice for robust cosmological emulation within the explored parameter space.

### 3.8. recommendations and future work

Based on these findings, the 1D CNN architecture is recommended for emulating the  $\Lambda$ CDM matter power spectrum within the tested parameter ranges due to its superior accuracy. Future work could explore hybrid architectures combining convolutional and dense layers, or incorporating attention mechanisms, to potentially further reduce outlier errors and improve robustness, especially in the high- $k$ , high- $z$  regime. Methods for uncertainty quantification (e.g., using Bayesian neural networks or ensemble methods) would enhance the utility of these emulators for precision cosmology. Targeted sampling or weighted loss functions could be employed during training to improve accuracy in specific regions of interest (e.g., BAO scales) or challenging parameter space corners. Finally, the methodology can be extended to emulate power spectra in more complex cosmological models by expanding the input parameter space and retraining the emulators.

## 4. CONCLUSIONS

In this study, we benchmarked two distinct neural network architectures, a 1D Convolutional Neural Network (CNN) and a Fully Connected Neural Network (FCNN), for the task of emulating the non-linear  $\Lambda$ CDM matter power spectrum  $P(k, z)$ . The goal was to identify an architecture that balances high accuracy with rapid inference speed, crucial for modern cosmological parameter inference pipelines.

We generated a large dataset of 50,000 power spectra using the Boltzmann code `classy_sz`, systematically sampling a 6-dimensional parameter space encompassing five standard  $\Lambda$ CDM parameters ( $\Omega_b h^2$ ,  $\Omega_{cdm} h^2$ ,  $H_0$ ,  $\log A$ ,  $n_s$ ) and redshift  $z$ , using Latin Hypercube Sampling. The power spectra were computed over a

fixed  $k$ -grid in the range  $[10^{-4}, 10]$  Mpc $^{-1}$  and processed in log-log space, with both input parameters and output  $\log_{10} P(k)$  values normalized for stable training. The data generation, network training, and performance evaluation were conducted as distinct computational steps.

Both the 1D CNN and FCNN models were trained on 45,000 samples and evaluated on an independent test set of 5,000 samples. The FCNN employed a standard feed-forward architecture, while the 1D CNN incorporated an initial dense projection followed by 1D convolutional layers to potentially leverage the sequential nature of the  $k$ -modes.

Our results demonstrate that the 1D CNN significantly outperforms the FCNN in terms of emulation accuracy. The 1D CNN achieved a mean absolute error of 0.0206 in  $\log_{10} P(k)$  on the test set, compared to 0.0313 for the FCNN. More importantly, the 95th percentile absolute error for the 1D CNN was 0.0751, substantially lower than the FCNN's 0.1210, indicating greater reliability across the parameter space. While both models exhibited rare large errors, the 1D CNN's maximum error was also lower. Visual inspection of emulated spectra confirmed that the 1D CNN provided a closer match to the true `classy_sz` spectra, particularly in the non-linear regime and at higher redshifts where the power spectrum shape is more complex.

Regarding computational speed, both emulators showed remarkably similar inference times when evaluated in batches (approximately 0.53 seconds for the entire 5,000-sample test set). This suggests that despite the 1D CNN having a significantly larger number of parameters (over 8 million vs. 1 million for the FCNN) and a larger memory footprint (32 MB vs. 4 MB), its architecture allows for efficient parallel processing on modern hardware, negating the potential speed disadvantage of its size for batch inference.

From these results, we learned that incorporating convolutional layers, even in a seemingly non-sequential input-to-output mapping like this, can be highly beneficial for emulating functions with inherent local structure and smoothness, such as the matter power spectrum in log-log space. The 1D CNN's ability to learn features across adjacent  $k$ -modes likely contributes to its superior performance, especially in capturing the details of the non-linear regime. The FCNN, treating the output vector as unstructured, is less effective at modeling these scale-dependent correlations.

The practical implications are significant. Both emulators provide orders-of-magnitude speedups compared to direct Boltzmann code evaluations, enabling rapid exploration of cosmological parameter space. However, the superior accuracy and robustness of the 1D CNN make it the preferred choice for precision cosmological analyses, such as likelihood calculations in MCMC or nested sampling frameworks, where reliable predictions across a wide range of parameters are critical. The modest increase in memory usage for the 1D CNN is generally acceptable for typical computing environments.

Future work could explore hybrid CNN-FCNN architectures, investigate methods for quantifying emulator uncertainty, or employ targeted sampling strategies to further reduce the frequency of outliers and improve accuracy in specific regions of parameter space (e.g., high  $z$ , extreme parameter values). Extending this methodology to emulate other cosmological observables or to models beyond  $\Lambda$ CDM would also be valuable.

In conclusion, our benchmarking study clearly demonstrates the advantage of using a 1D Convolutional Neural Network architecture over a standard Fully Connected Neural Network for emulating the non-linear  $\Lambda$ CDM matter power spectrum. The 1D CNN provides significantly higher accuracy and greater robustness across the sampled parameter and redshift range, with comparable inference speed, making it a highly effective tool for accelerating cosmological analyses.

## REFERENCES

- Adams, A. D., Colose, C., Merrelli, A., Turnbull, M., & Kane, S. R. 2024, Habitability in 4-D: Predicting the Climates of Earth Analogs across Rotation and Orbital Configurations. <https://arxiv.org/abs/2412.19357>
- Andati, L. A. L., Smirnov, O. M., Makhathini, S., & Sebokolodi, L. M. 2023, PolarVis: Towards Web-based Polarimetric Analysis. <https://arxiv.org/abs/2312.07645>
- Bartlett, D. J., Kammerer, L., Kronberger, G., et al. 2024, A precise symbolic emulator of the linear matter power spectrum, doi: <https://doi.org/10.1051/0004-6361/202348811>
- Bevins, H. T. J., Gessey-Jones, T., & Handley, W. J. 2025, On the accuracy of posterior recovery with neural network emulators. <https://arxiv.org/abs/2503.13263>

- Bolliet, B., Mancini, A. S., Hill, J. C., et al. 2023, High-accuracy emulators for observables in  $\Lambda$ CDM,  $N_{\text{eff}}$ ,  $\Sigma m_\nu$ , and  $w$  cosmologies. <https://arxiv.org/abs/2303.01591>
- Chabanier, S., Millea, M., & Palanque-Delabrouille, N. 2019, Matter power spectrum: from Ly $\alpha$  forest to CMB scales, doi: <https://doi.org/10.1093/mnras/stz2310>
- Chandra, K. 2024, gwforge: A user-friendly package to generate gravitational-wave mock data. <https://arxiv.org/abs/2407.21109>
- Collaboration, P., Ade, P. A. R., Aghanim, N., et al. 2016, Planck 2015 results. XIII. Cosmological parameters, doi: <https://doi.org/10.1051/0004-6361/201525830>
- Collaboration, P., Aghanim, N., Akrami, Y., et al. 2021, Planck 2018 results. VI. Cosmological parameters, doi: <https://doi.org/10.1051/0004-6361/201833910>
- Contardo, G., Trotta, R., Gioia, S. D., Hogg, D. W., & Villaescusa-Navarro, F. 2025, On the effects of parameters on galaxy properties in CAMELS and the predictability of  $\Omega_m$ . <https://arxiv.org/abs/2503.22654>
- Cowell, J. A., Alonso, D., & Liu, J. 2024, Hitting the mark: Optimising Marked Power Spectra for Cosmology, doi: <https://doi.org/10.1093/mnras/stae2492>
- Daoutis, C., Zezas, A., Kyritsis, E., Kouroumpatzakis, K., & Bonfini, P. 2024, From seagull to hummingbird: New diagnostic methods for resolving galaxy activity. <https://arxiv.org/abs/2411.08983>
- Dekel, A., Burstein, D., & White, S. D. M. 1996, Measuring Omega. <https://arxiv.org/abs/astro-ph/9611108>
- DeRose, J., Chen, S.-F., White, M., & Kokron, N. 2021, Neural Network Acceleration of Large-scale Structure Theory Calculations, doi: <https://doi.org/10.1088/1475-7516/2022/04/056>
- Dorn-Wallenstein, T. Z., & Levesque, E. M. 2018, Stellar Population Diagnostics of the Massive Star Binary Fraction, doi: <https://doi.org/10.3847/1538-4357/aae5d6>
- Fisher, C., & Heng, K. 2022, How do we optimally sample model grids of exoplanet spectra?, doi: <https://doi.org/10.3847/1538-4357/ac7801>
- Fluri, J., Lucchi, A., Kacprzak, T., Refregier, A., & Hofmann, T. 2021, Cosmological Parameter Estimation and Inference using Deep Summaries, doi: <https://doi.org/10.1103/PhysRevD.104.123526>
- Gebran, M. 2024, Generating stellar spectra using Neural Networks. <https://arxiv.org/abs/2401.13411>
- Günther, S., Balkenhol, L., Fidler, C., et al. 2025, OLÉ – Online Learning Emulation in Cosmology. <https://arxiv.org/abs/2503.13183>
- Jamieson, D., Li, Y., de Oliveira, R. A., et al. 2022, Field Level Neural Network Emulator for Cosmological N-body Simulations, doi: <https://doi.org/10.3847/1538-4357/acdb6c>
- Jamieson, D., Li, Y., Villaescusa-Navarro, F., Ho, S., & Spergel, D. N. 2024, Field-level Emulation of Cosmic Structure Formation with Cosmology and Redshift Dependence. <https://arxiv.org/abs/2408.07699>
- Jense, H. T., Harrison, I., Calabrese, E., et al. 2024, A complete framework for cosmological emulation and inference with CosmoPower. <https://arxiv.org/abs/2405.07903>
- Lehman, K., Schuster, N., Lucie-Smith, L., et al. 2025, Cosmological Inference with Cosmic Voids and Neural Network Emulators. <https://arxiv.org/abs/2502.05262>
- Marulanda, J. P., Santa, C., & Romano, A. E. 2020, Deep learning merger masses estimation from gravitational waves signals in the frequency domain, doi: <https://doi.org/10.1016/j.physletb.2020.135790>
- Mathews, E. P., Leja, J., Speagle, J. S., et al. 2023, As Simple as Possible but No Simpler: Optimizing the Performance of Neural Net Emulators for Galaxy SED Fitting. <https://arxiv.org/abs/2306.16442>
- Min, Z., Xiao, X., Ding, J., et al. 2024, Deep learning for cosmological parameter inference from a dark matter halo density field, doi: <https://doi.org/10.1103/PhysRevD.110.063531>
- Monsalves, N., Arancibia, M. J., Bayo, A., et al. 2024, Application of Convolutional Neural Networks to time domain astrophysics. 2D image analysis of OGLE light curves, doi: <https://doi.org/10.1051/0004-6361/202449995>
- Narkedimilli, S., Amballa, V. S., Kumar, N. V. S., et al. 2025, Comparative Analysis of Black Hole Mass Estimation in Type-2 AGNs: Classical vs. Quantum Machine Learning and Deep Learning Approaches. <https://arxiv.org/abs/2502.15297>
- Niu, J., He, P., & Zhang, T.-J. 2025, Constraining the Hubble Constant with a Simulated Full Covariance Matrix Using Neural Networks. <https://arxiv.org/abs/2502.11443>
- Ocampo, I., Cañas-Herrera, G., & Nesseris, S. 2025, Neural Networks for cosmological model selection and feature importance using Cosmic Microwave Background data, doi: <https://doi.org/10.1088/1475-7516/2025/02/004>
- Orjuela-Quintana, J. B., Nesseris, S., & Sapone, D. 2024a, Machine learning unveils the linear matter power spectrum of modified gravity, doi: <https://doi.org/10.1103/PhysRevD.109.063511>

- Orjuela-Quintana, J. B., Sapone, D., & Nesseris, S. 2024b, Analytical Emulator for the Linear Matter Power Spectrum from Physics-Informed Machine Learning. <https://arxiv.org/abs/2407.16640>
- Oxholm, T. M. 2022, A Beginner's Guide to Line Intensity Mapping Power Spectra, doi: <https://doi.org/10.17307/wsc.v1i1.352>
- Padilla, L. E., Tellez, L. O., Escamilla, L. A., & Vazquez, J. A. 2021, Cosmological parameter inference with Bayesian statistics, doi: <https://doi.org/10.3390/universe7070213>
- Pan, S., Liu, M., Forero-Romero, J., et al. 2020, Cosmological parameter estimation from large-scale structure deep learning. <https://arxiv.org/abs/1908.10590>
- Preston, C., Amon, A., & Efstathiou, G. 2024, Reconstructing the matter power spectrum with future cosmic shear surveys. <https://arxiv.org/abs/2404.18240>
- Raghav, S., Ayitapu, P., Narkedimilli, S., Makam, S., & H. A. B. 2024, Photometric Analysis for Predicting Star Formation Rates in Large Galaxies Using Machine Learning and Deep Learning Techniques. <https://arxiv.org/abs/2410.06736>
- Rampf, C., Schobesberger, S. O., & Hahn, O. 2022, Analytical growth functions for cosmic structures in a  $\Lambda$ CDM Universe, doi: <https://doi.org/10.1093/mnras/stac2406>
- Reza, M., Zhang, Y., Nord, B., et al. 2022, Estimating Cosmological Constraints from Galaxy Cluster Abundance using Simulation-Based Inference. <https://arxiv.org/abs/2208.00134>
- Różański, T., Ting, Y.-S., & Jabłońska, M. 2024, TransformerPayne: enhancing spectral emulation accuracy and data efficiency by capturing long-range correlations. <https://arxiv.org/abs/2407.05751>
- Saraivanov, E., Zhong, K., Miranda, V., et al. 2024, Attention-Based Neural Network Emulators for Multi-Probe Data Vectors Part II: Assessing Tension Metrics. <https://arxiv.org/abs/2403.12337>
- Silva, E., Zúñiga-Bolaño, U., Nunes, R. C., & Valentino, E. D. 2024, Non-Linear Matter Power Spectrum Modeling in Interacting Dark Energy Cosmologies, doi: <https://doi.org/10.1140/epjc/s10052-024-13487-x>
- Smith, M. J., & Geach, J. E. 2023, Astronomia ex machina: a history, primer, and outlook on neural networks in astronomy, doi: <https://doi.org/10.1098/rsos.221454>
- Tanigawa, S., Glazebrook, K., Jacobs, C., Labbe, I., & Qin, A. K. 2024, HAYATE: Photometric redshift estimation by hybridising machine learning with template fitting. <https://arxiv.org/abs/2402.00323>
- Thob, A. C. R., Sanderson, R. E., Eden, A. P., et al. 2024, Generating synthetic star catalogs from simulated data for next-gen observatories with py-ananke, doi: <https://doi.org/10.21105/joss.06234>
- Thomas, D. B., Kopp, M., & Skordis, C. 2016, Constraining dark matter properties with Cosmic Microwave Background observations, doi: <https://doi.org/10.3847/0004-637X/830/2/155>
- Tripathi, A., Kaur, G., Datta, A., & Majumdar, S. 2024, Comparing sampling techniques to chart parameter space of 21 cm Global signal with Artificial Neural Networks, doi: <https://doi.org/10.1088/1475-7516/2024/10/041>
- Trusov, S., Zarrouk, P., & Cole, S. 2025, Neural Network-based model of galaxy power spectrum: Fast full-shape galaxy power spectrum analysis. <https://arxiv.org/abs/2403.20093>
- Urbán, J. F., Stefanou, P., & Pons, J. A. 2024, Unveiling the optimization process of Physics Informed Neural Networks: How accurate and competitive can PINNs be? <https://arxiv.org/abs/2405.04230>
- Veiga, M. H., Meng, X., Gnedin, O. Y., Gnedin, N. Y., & Huan, X. 2021, Reconstruction of the Density Power Spectrum from Quasar Spectra using Machine Learning. <https://arxiv.org/abs/2107.09082>